

DTU



02/11/23

Real-Time Visual and Machine Learning Systems

Questions or comments from last time?

Agenda – Module 1 part A

- Memory Hierarchies in Hardware
- Memory Hierarchies in Software
- Memory Allocations and Data Structures
- Smart Pointers
- Graph Structures
- Garbage Collectors
- Computational Graphs
- Exercises

Memory Hierarchies in Hardware

- The speed of RAM and disks vs. the speed of processors
- Not the whole picture – core count, cache sizes, efficiency, sophistication of branch prediction, bus bandwidth

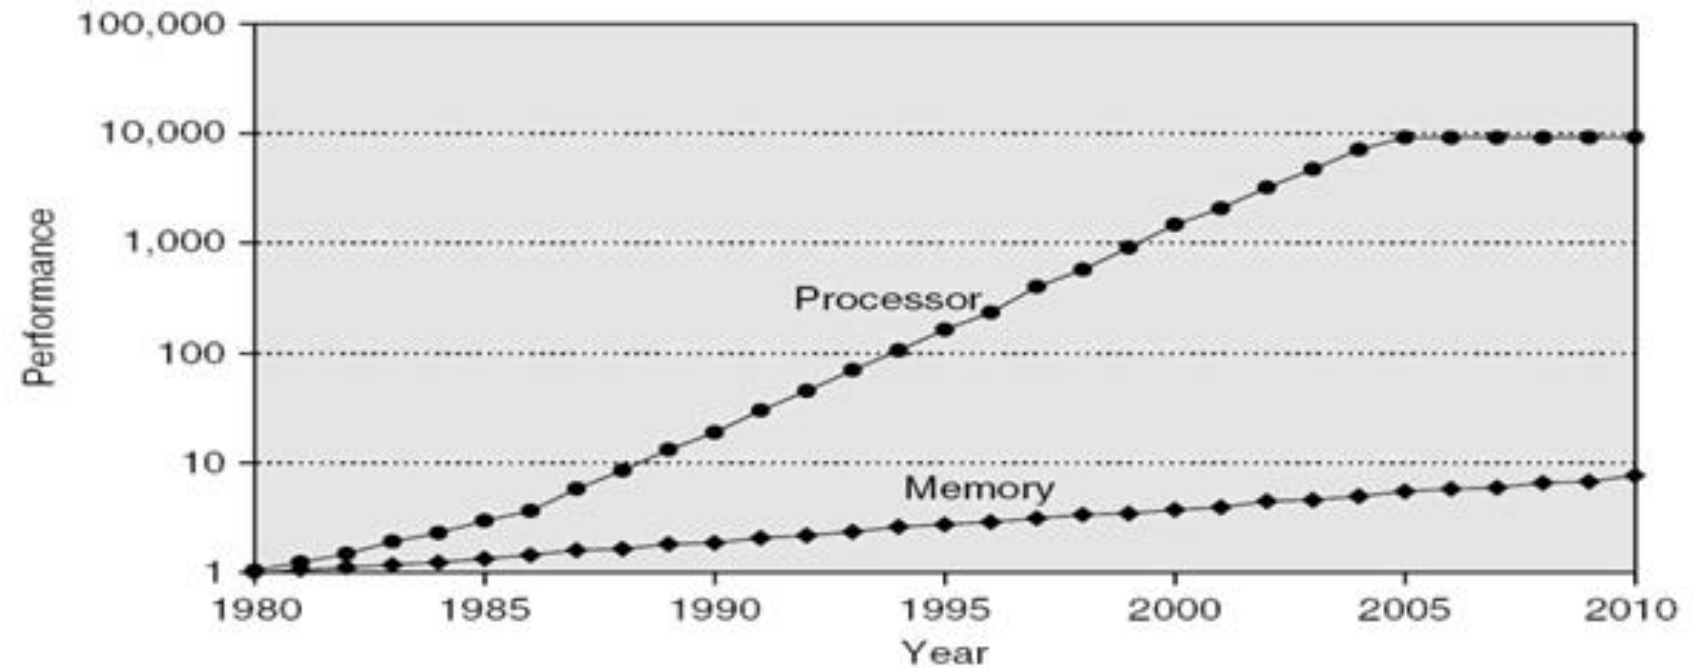


Figure 25.2

Memory Hierarchies in Hardware

From slowest to fastest

CPU Hierarchy Simplified

Faster/Smaller



Slower/Larger

Registers

L1 Cache

L2 Cache

L3 Cache

Memory (RAM)

Storage (HDD, SSD, NVMe)

Memory Hierarchies in Hardware

Different levels can be shared differently

2-Core CPU Hierarchy

Faster/Smaller



Slower/Larger

CPU0 - Registers

CPU1 - Registers

CPU0 -
L1 Cache

CPU1 -
L1 Cache

CPU 0 & 1 - L2 Cache

L3 Cache

Memory (RAM)

Storage (HDD, SSD, NVMe)

Memory Hierarchies in Hardware

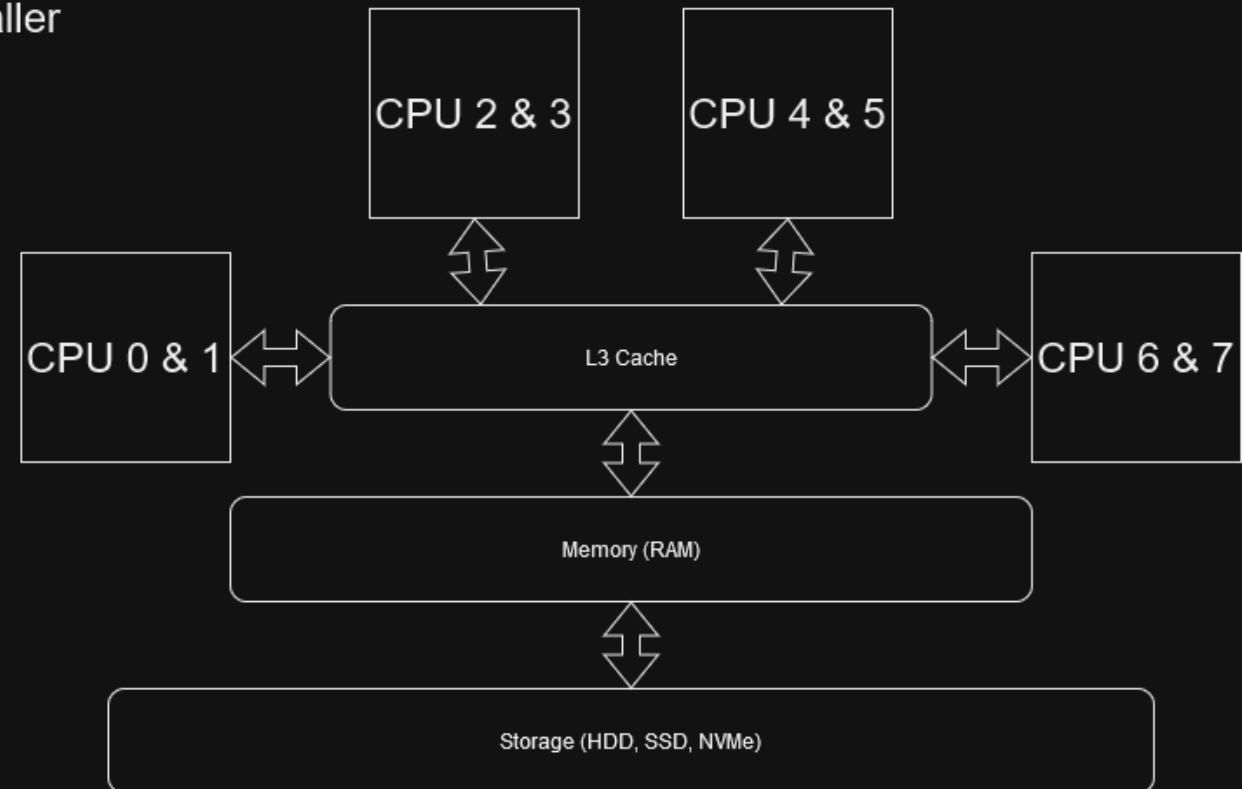
Different levels can be shared differently

8-Core CPU Hierarchy

Faster/Smaller



Slower/Larger



Memory Hierarchies in Hardware

We can't manipulate the caches directly

```
for row_output in 0..output.row_count {  
  for column_output in 0..output.column_count {  
    for inner_dimension in 0..input.column_count {  
      output.data[row_output * output.column_count + column_output] += input.data  
        [row_output * input.column_count + inner_dimension]  
        * weights.data[inner_dimension * weights.column_count + column_output];  
    }  
  }  
}
```

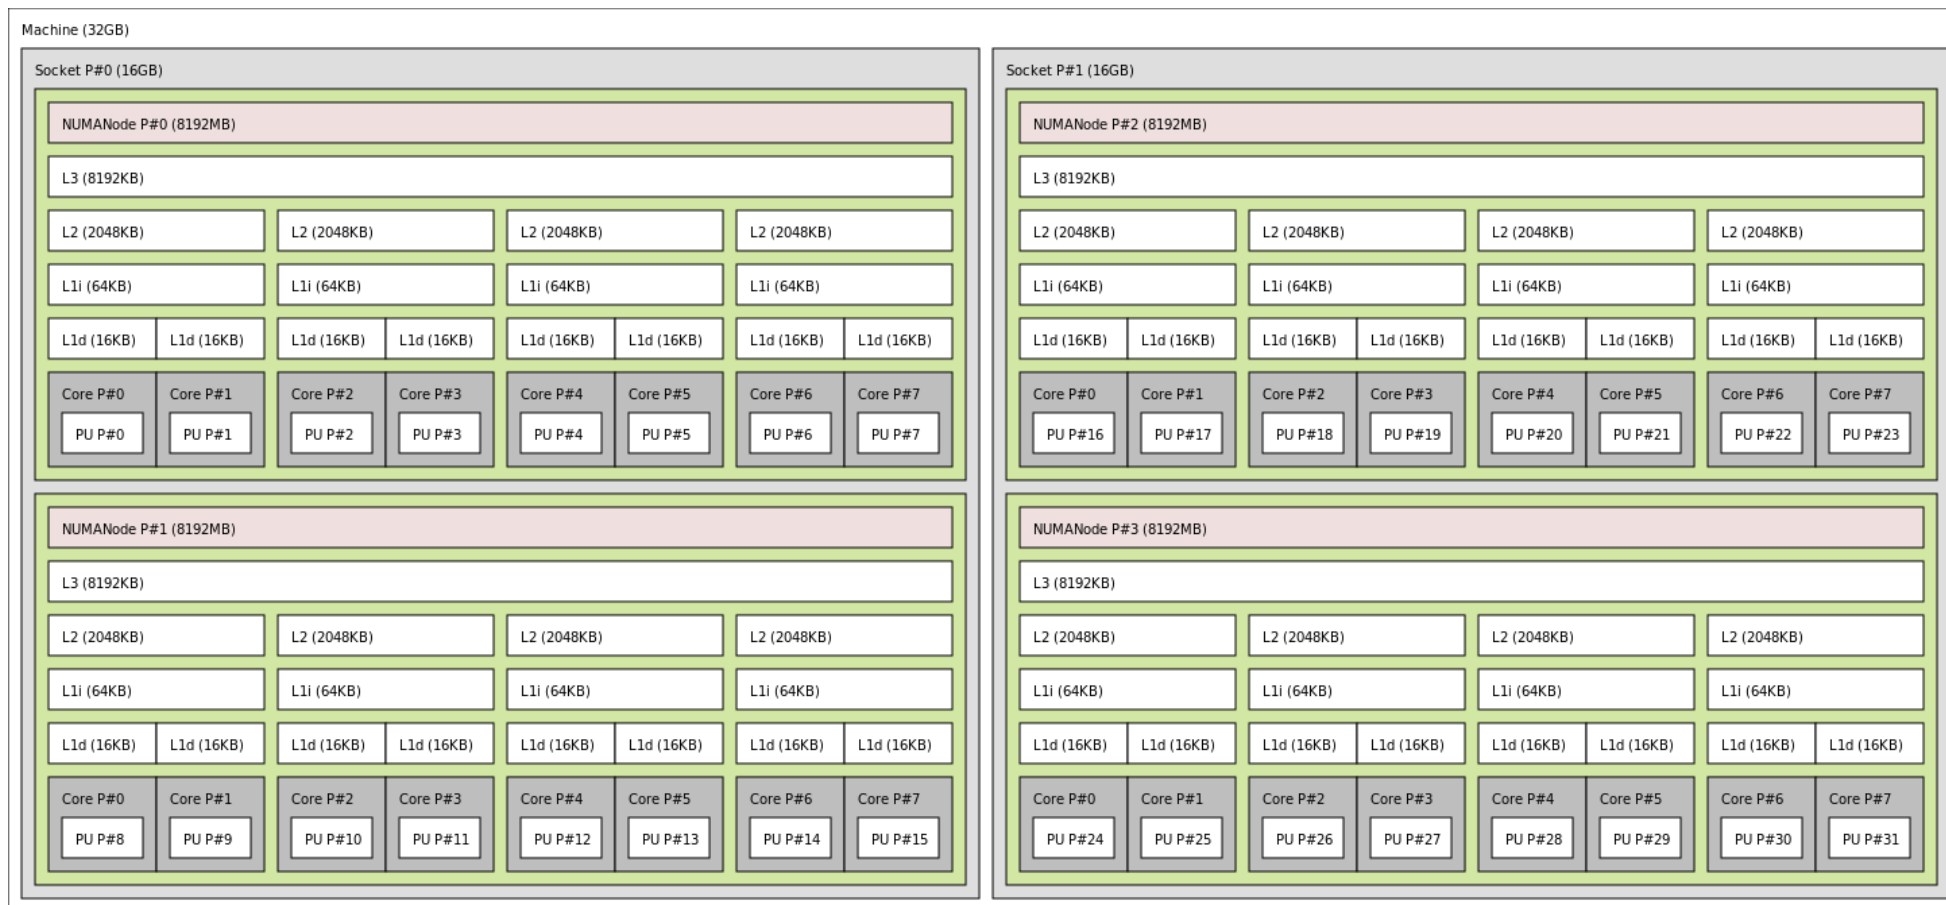
Should have cleaned up the linearized
accesses

Memory Hierarchies in Hardware

But we can create stack variables of known sizes, which are likely to reside in registers and L1 cache

```
for row_output in 0..output.row_count {  
  for column_output in 0..output.column_count {  
    let mut result: f32 = 0.0;  
    for inner_dimension in 0..input.column_count {  
      result += input.data[row_output * input.column_count + inner_dimension]  
        * weights.data[inner_dimension * weights.column_count + column_output];  
    }  
    output.data[row_output * output.column_count + column_output] = result;  
  }  
}
```

Memory Hierarchies in Hardware – AMD Bulldozer server



[Image Link](#)

Memory Hierarchies in Hardware – GPU (H100)

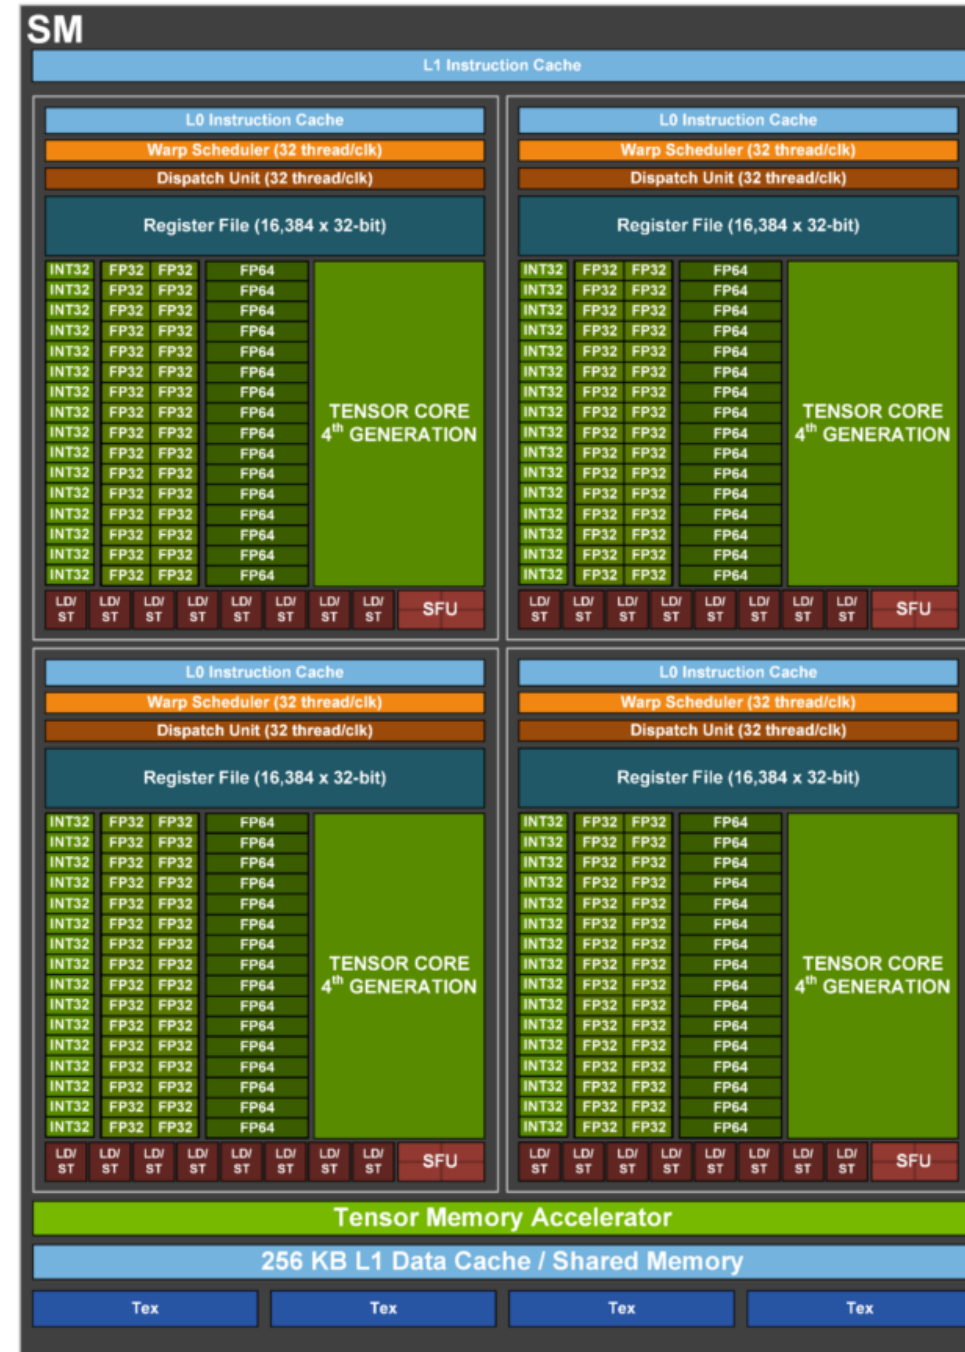


[Image Link](#)

Memory Hierarchies in Hardware GPU (H100)

More about GPU architecture next week

[Image Link](#)



Memory Hierarchies in Software

Memory Allocations – in C!

```
int element_count = 42;  
int* base_integer_array = malloc(element_count * sizeof(int));  
  
*base_integer_array = 0;  
*(base_integer_array + 1) = 1;  
base_integer_array[2] = 2;  
  
int* integer_array = base_integer_array + 3;  
*integer_array = 3;  
integer_array[1] = 4;
```


Memory Hierarchies in Software

Memory Allocations – Spot the Error!

```
int element_count = 42;  
int* base_integer_array = malloc(element_count * sizeof(int));  
  
*base_integer_array = 0;  
*(base_integer_array + 1) = 1;  
base_integer_array[2] = 2;  
  
int* integer_array = base_integer_array + 3;  
*integer_array = 3;  
integer_array[1] = 4;  
  
free(integer_array);
```

Memory Hierarchies in Software

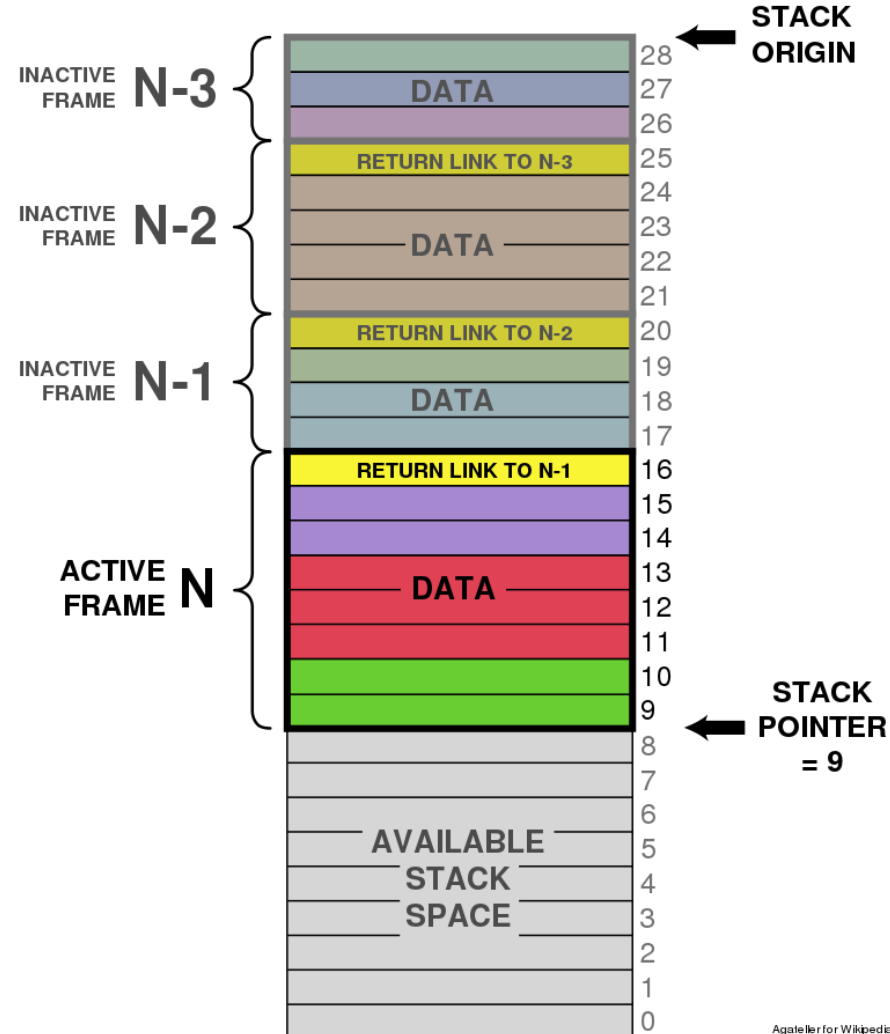
Memory Allocations – Spot the Error!

```
int element_count = 42;  
int* base_integer_array = malloc(element_count * sizeof(int));  
  
*base_integer_array = 0;  
*(base_integer_array + 1) = 1;  
base_integer_array[2] = 2;  
  
int* integer_array = base_integer_array + 3;  
*integer_array = 3;  
integer_array[1] = 4;  
  
free(base_integer_array);
```


Memory Hierarchies in Software

Where does the memory come from?

- Function calls
- Variables with a size known at compile time




[Image Link](#)

Memory Hierarchies in Software

Where does the memory come from?

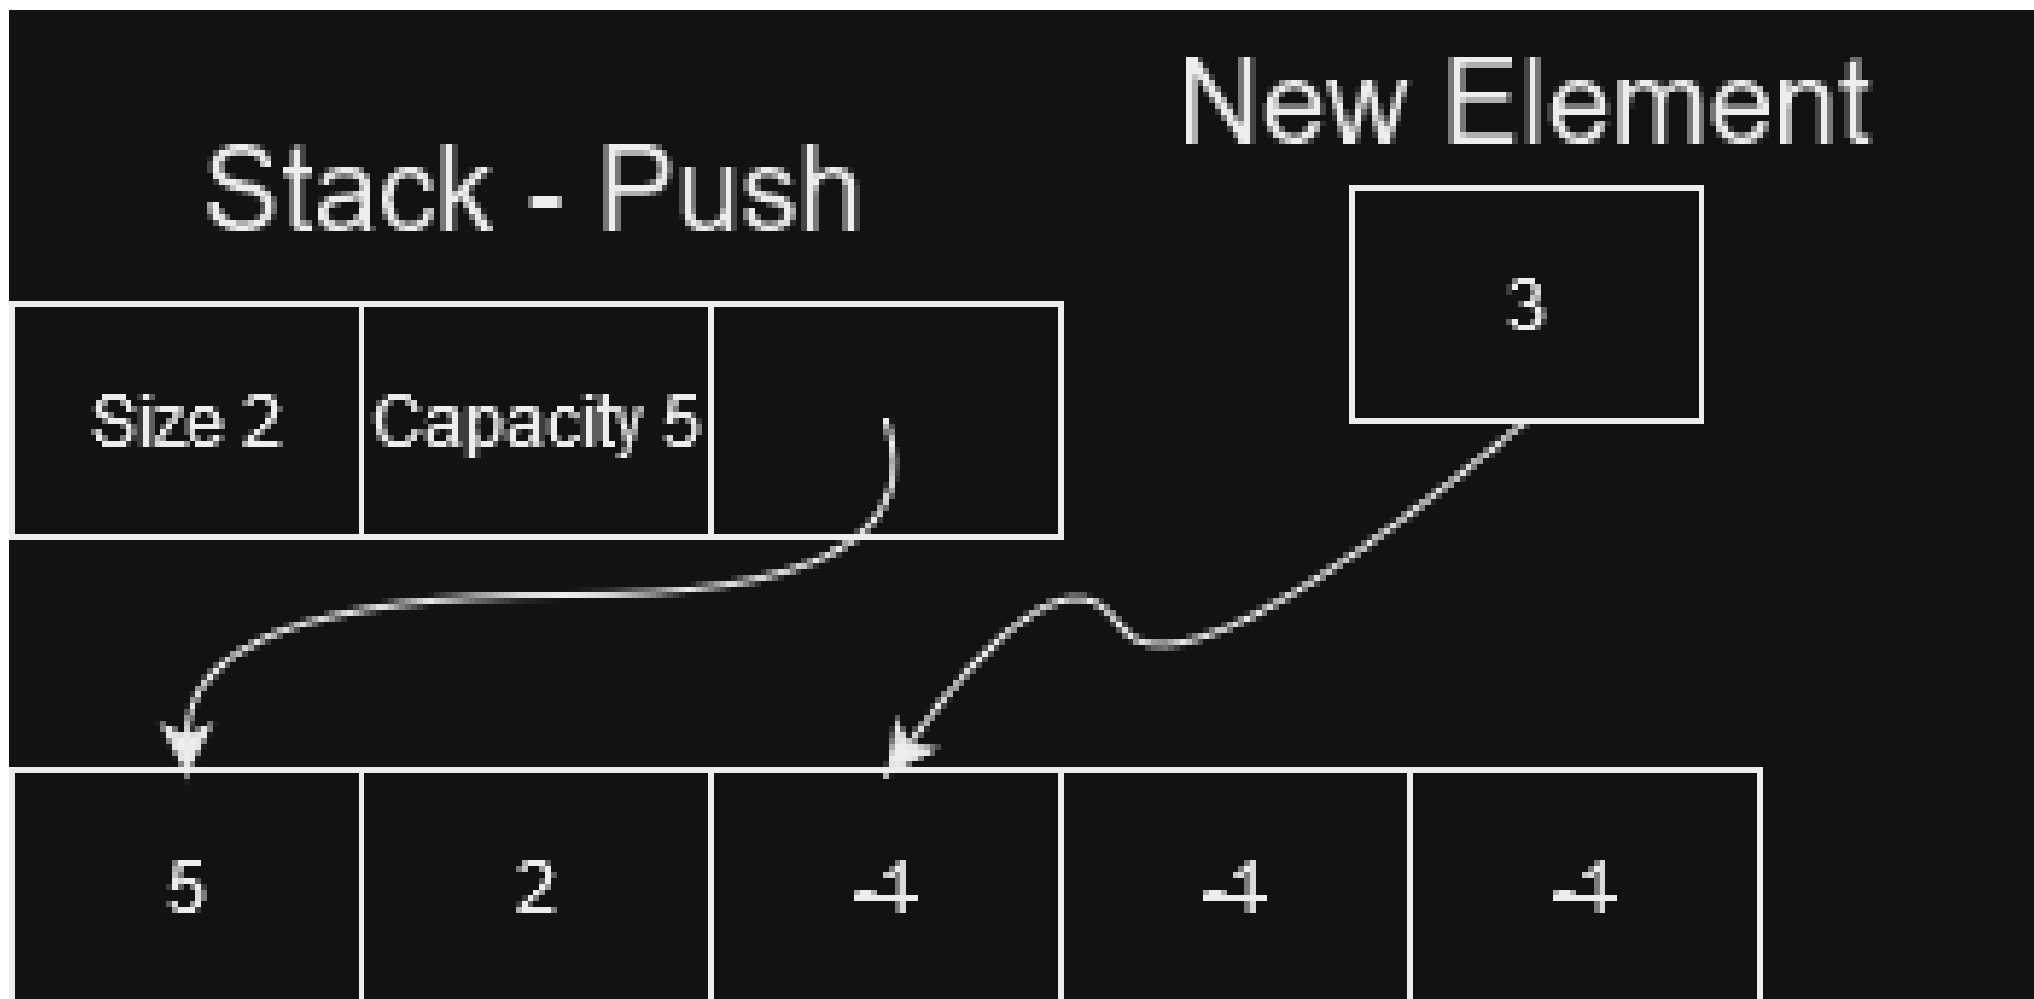
- When we malloc'ed, the data was allocated on the heap



[Image Link](#)

Memory Allocations and Data Structures

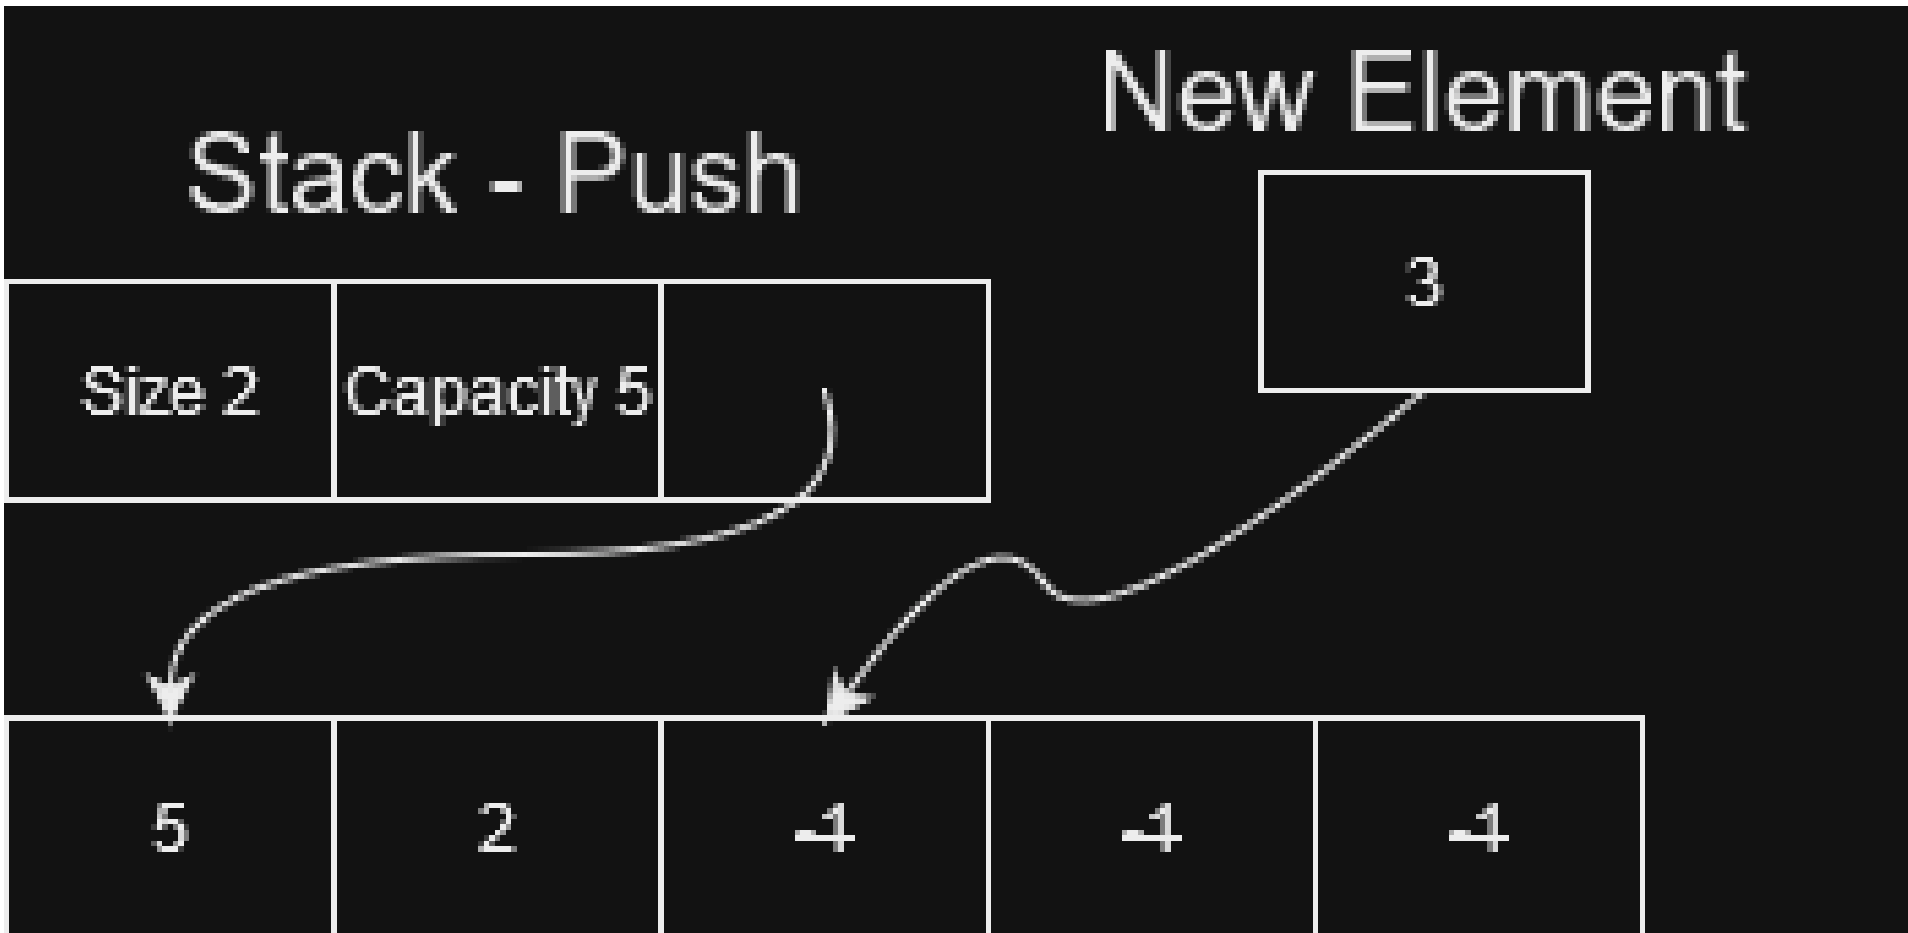
Dynamic Arrays



Memory Allocations and Data Structures

Dynamic Arrays

Why resizing can be problematic for memory safety (pointers and references).



Memory Hierarchies in Software


Reference, Copy, Clone, Move

Adding references is the default in Python

Clone is the default in C++ for most types like `vector<float>`

Move is the default in Rust for most types like `Vec<f32>`

Copy is standard for very small types like `f32` and `u32` for both languages



Memory Allocations and Data Structures Accesses and Strides

Now that we know how a
list/vector/dynamic array works,
how do we access it?

Strides

```
RUNNING ACCESS TESTS WITH 100 data elements for 100000 iterations!  
=====
```

Sequential access:	1 ms
Non-wrapping strided access (2):	1 ms
Non-wrapping strided access (3):	0 ms
Non-wrapping strided access (4):	0 ms
Strided access (1):	1 ms
Strided access (5):	15 ms
Strided access (17):	14 ms
Random access:	77 ms

```
RUNNING ACCESS TESTS WITH 1000 data elements for 100000 iterations!  
=====
```

Sequential access:	8 ms
Non-wrapping strided access (2):	13 ms
Non-wrapping strided access (3):	11 ms
Non-wrapping strided access (4):	6 ms
Strided access (1):	9 ms
Strided access (5):	134 ms
Strided access (17):	139 ms
Random access:	472 ms

```
RUNNING ACCESS TESTS WITH 10000 data elements for 100000 iterations!  
=====
```

Sequential access:	89 ms
Non-wrapping strided access (2):	125 ms
Non-wrapping strided access (3):	86 ms
Non-wrapping strided access (4):	63 ms
Strided access (1):	94 ms
Strided access (5):	1311 ms
Strided access (17):	1328 ms
Random access:	19206 ms

```
RUNNING ACCESS TESTS WITH 100000 data elements for 100000 iterations!  
=====
```

Sequential access:	960 ms
Non-wrapping strided access (2):	1510 ms
Non-wrapping strided access (3):	1081 ms
Non-wrapping strided access (4):	787 ms
Strided access (1):	1003 ms
Strided access (5):	14595 ms
Strided access (17):	15034 ms
Random access:	86041 ms

Memory Allocations and Data Structures

Accesses and Strides

Cache Lines

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
a.first	a.first	a.second	a.second	a.second	a.second	b.first	b.first	b.second	b.second	b.second	b.second	c.first	c.first	c.second	c.second

Memory Hierarchies in Hardware

From slowest to fastest

CPU Hierarchy Simplified

Faster/Smaller



Slower/Larger

Registers

L1 Cache

L2 Cache

L3 Cache

Memory (RAM)

Storage (HDD, SSD, NVMe)

Memory Allocations and Data Structures

Accesses and Strides

Pad for alignment

Remove pad for storage

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
a.first	a.first	a.pad	a.pad	a.second	a.second	a.second	a.second	b.first	b.first	b.pad	b.pad	b.second	b.second	c.second	c.second

Memory Allocations and Data Structures Accesses and Strides

Strides – are you getting the most out of your cache line?

RUNNING ACCESS TESTS WITH 100 data elements for 100000 iterations!

```
=====
Sequential access: 1 ms
Non-wrapping strided access (2): 1 ms
Non-wrapping strided access (3): 0 ms
Non-wrapping strided access (4): 0 ms
Strided access (1): 1 ms
Strided access (5): 15 ms
Strided access (17): 14 ms
Random access: 77 ms
```

RUNNING ACCESS TESTS WITH 1000 data elements for 100000 iterations!

```
=====
Sequential access: 8 ms
Non-wrapping strided access (2): 13 ms
Non-wrapping strided access (3): 11 ms
Non-wrapping strided access (4): 6 ms
Strided access (1): 9 ms
Strided access (5): 134 ms
Strided access (17): 139 ms
Random access: 472 ms
```

RUNNING ACCESS TESTS WITH 10000 data elements for 100000 iterations!

```
=====
Sequential access: 89 ms
Non-wrapping strided access (2): 125 ms
Non-wrapping strided access (3): 86 ms
Non-wrapping strided access (4): 63 ms
Strided access (1): 94 ms
Strided access (5): 1311 ms
Strided access (17): 1328 ms
Random access: 19206 ms
```

RUNNING ACCESS TESTS WITH 100000 data elements for 100000 iterations!

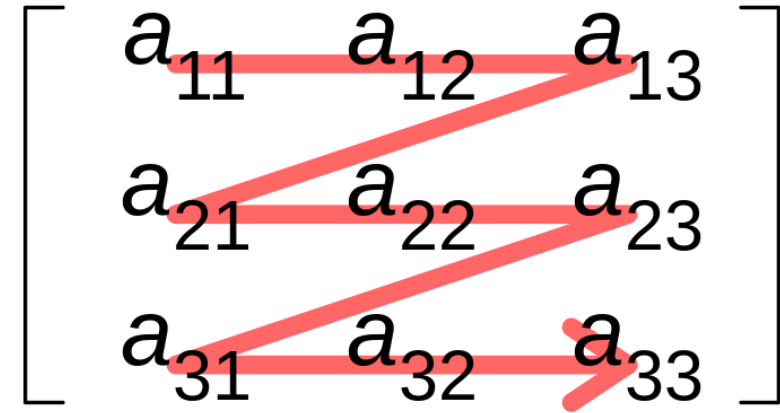
```
=====
Sequential access: 960 ms
Non-wrapping strided access (2): 1510 ms
Non-wrapping strided access (3): 1081 ms
Non-wrapping strided access (4): 787 ms
Strided access (1): 1003 ms
Strided access (5): 14595 ms
Strided access (17): 15034 ms
Random access: 86041 ms
```

Memory Allocations and Data Structures

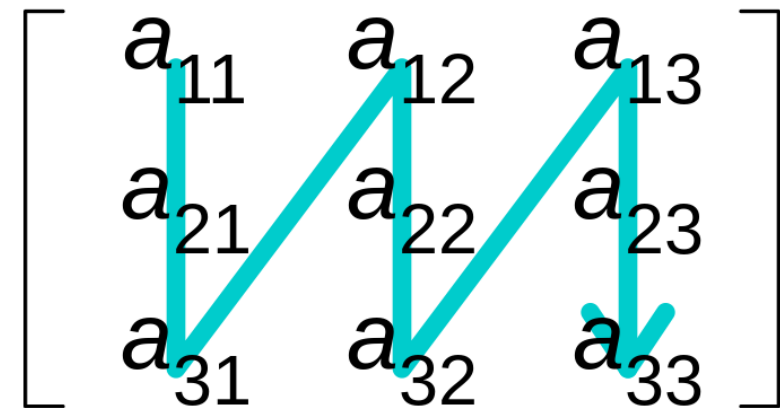
ND Arrays

Row-major and Column-major ordering

Row-major order



Column-major order



Memory Allocations and Data Structures

ND Arrays

We can do this in
2D with row major
ordering

```
let data: [[[i32; 2]; 2]; 2] =  
    [  
        [[1, 2], [3, 4]],  
        [[5, 6], [7, 8]]  
    ];  
  
let x_dimension: usize = 2;  
let y_dimension: usize = 2;  
let z_dimension: usize = 2;  
  
for x_index in 0..x_dimension {  
    for y_index in 0..y_dimension {  
        for z_index in 0..z_dimension {  
            println("{} ", data[x_index][y_index][z_index]);  
        }  
    }  
}
```

Memory Allocations and Data Structures

ND Arrays

We can do this in
2D with column major
ordering

```
let data: [[[i32; 2]; 2]; 2] =  
    [  
        [[1, 2], [3, 4],  
         [5, 6], [7, 8]]  
    ];  
  
let x_dimension: usize = 2;  
let y_dimension: usize = 2;  
let z_dimension: usize = 2;  
  
for z_index in 0..z_dimension {  
    for y_index in 0..y_dimension {  
        for x_index in 0..x_dimension {  
            println("{} ", data[x_index][y_index][z_index]);  
        }  
    }  
}
```

Memory Allocations and Data Structures

ND Arrays – Linearized Indexing

Or in 1D and gain nice properties

```
let mut data: Vec<i32> = Vec::<i32>::new();
data.push(vec![0, 1, 2, 3]);

let column_count: usize = 2;
let row_count: usize = 2;

for x_index in 0..row_count {
    for y_index in 0..column_count {
        println!("{}", data[x_index * column_count + y_index]);
    }
}
```

Memory Allocations and Data Structures

ND Arrays – Linearized Indexing

$$x_index * y_size * z_size + y_index * z_size + z_index$$


With linearized indexing we gain nice properties like different views on the same data, permutations, easier data transfers and better performance.

It also becomes a lot easier to move your code to the GPU.

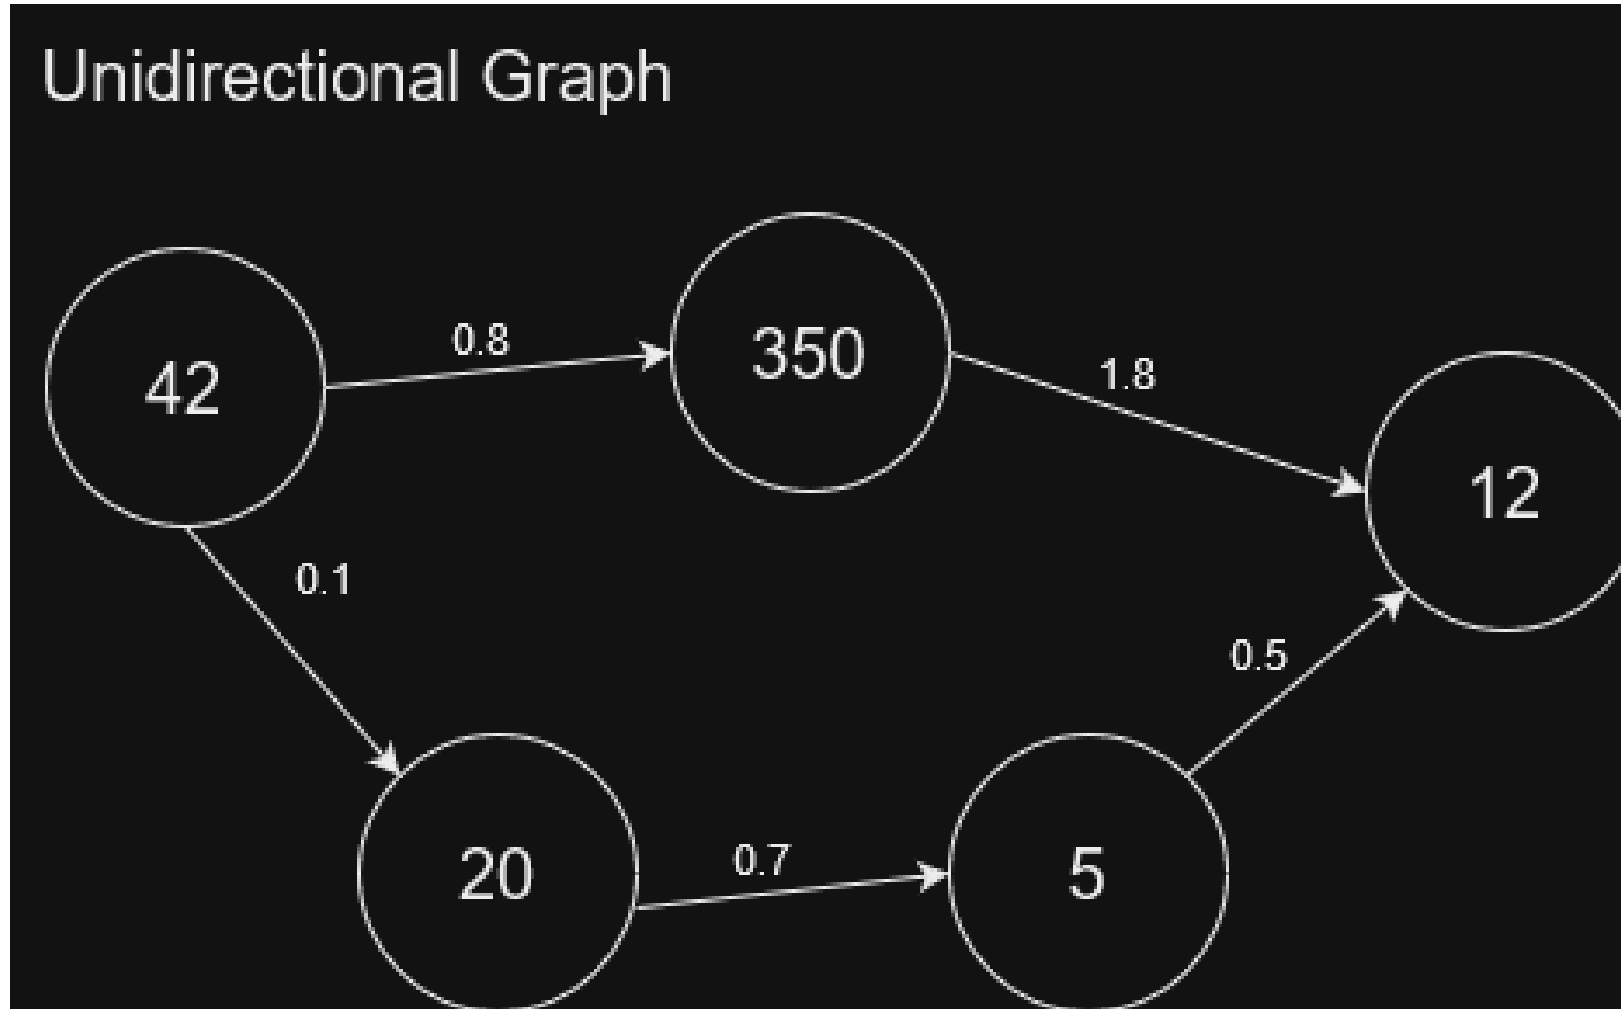
GRAPHS!

Pointers, pointers, pointers


Graph Structures - Bidirectional



Graph Structures – Unidirectional with weighted edges



Graph Structures - Trees

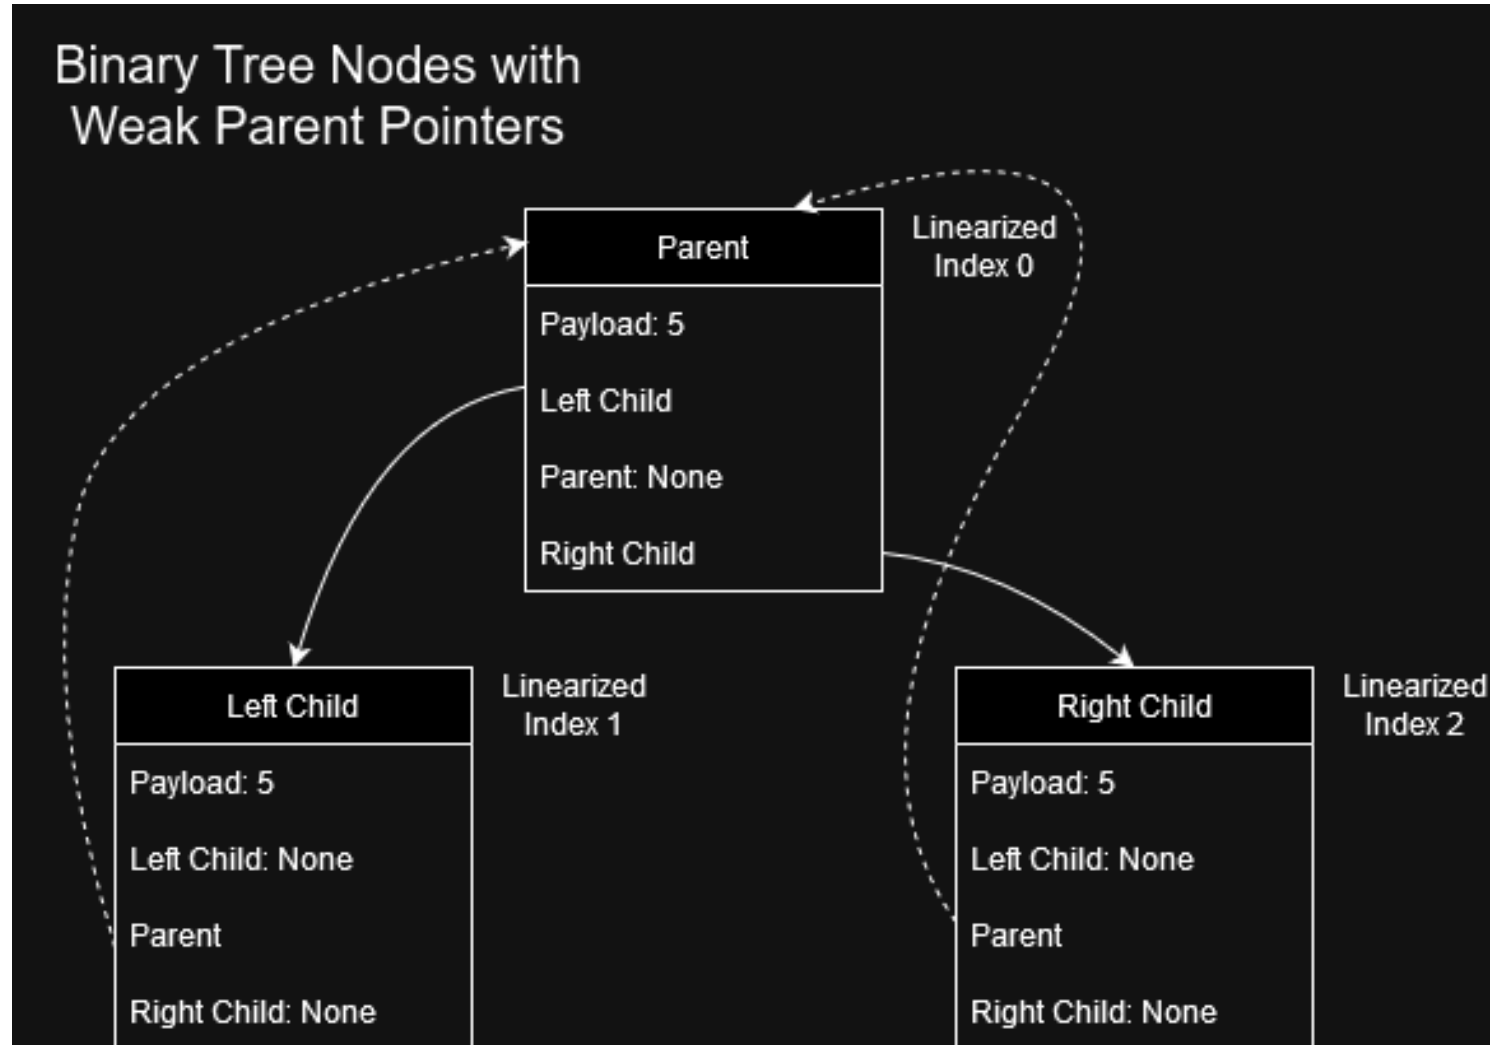


Smart Pointers

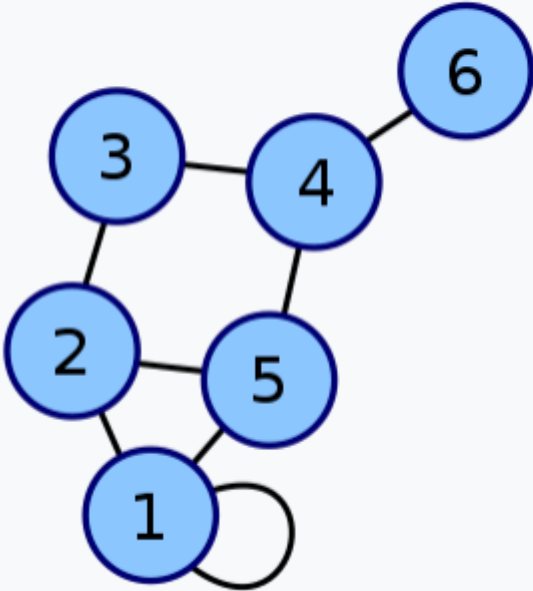
Rust and C++11+

- `Box<T>` or `unique_ptr<T>`
- `Rc<T>`, `Arc<T>` and `shared_ptr<T>`
- `Weak<T>` and `weak_ptr<T>`

Graph Structures



Graph Structures with Indices

Labeled graph	Adjacency matrix
	$\begin{pmatrix} 2 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$ <p>Coordinates are 1–6.</p>

[Image Link](#)

Safety and Pointer Based Data Structures

Such a hard topic that there is a small book written about how to even do this safely in Rust

[Learning Rust With Entirely Too Many Linked Lists](#)

Computational Graphs

```
# x represents our data
def forward(self, x):
    # Pass data through conv1
    x = self.conv1(x)
    # Use the rectified-linear activation function over x
    x = F.relu(x)

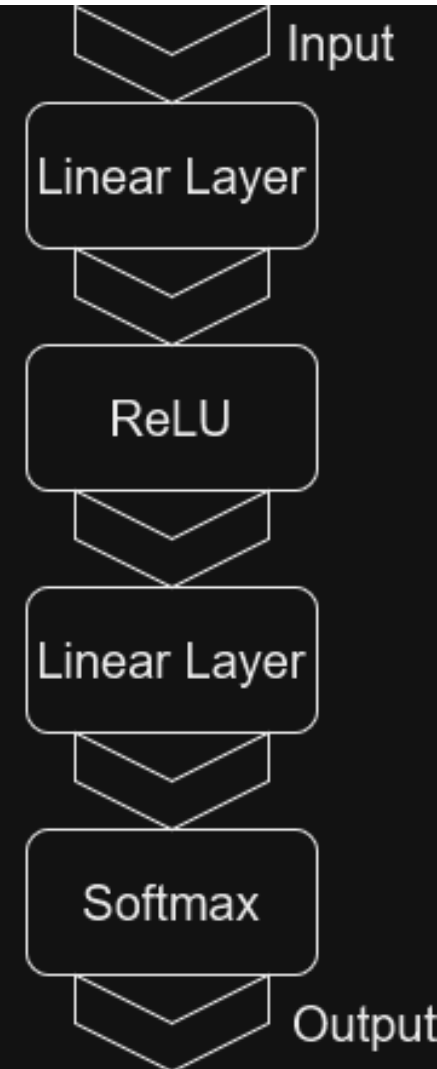
    x = self.conv2(x)
    x = F.relu(x)

    # Run max pooling over x
    x = F.max_pool2d(x, 2)
    # Pass data through dropout1
    x = self.dropout1(x)
    # Flatten x with start_dim=1
    x = torch.flatten(x, 1)
    # Pass data through ``fc1``
    x = self.fc1(x)
    x = F.relu(x)
    x = self.dropout2(x)
    x = self.fc2(x)


    # Apply softmax to x
    output = F.log_softmax(x, dim=1)
    return output
```

[Image Link](#)

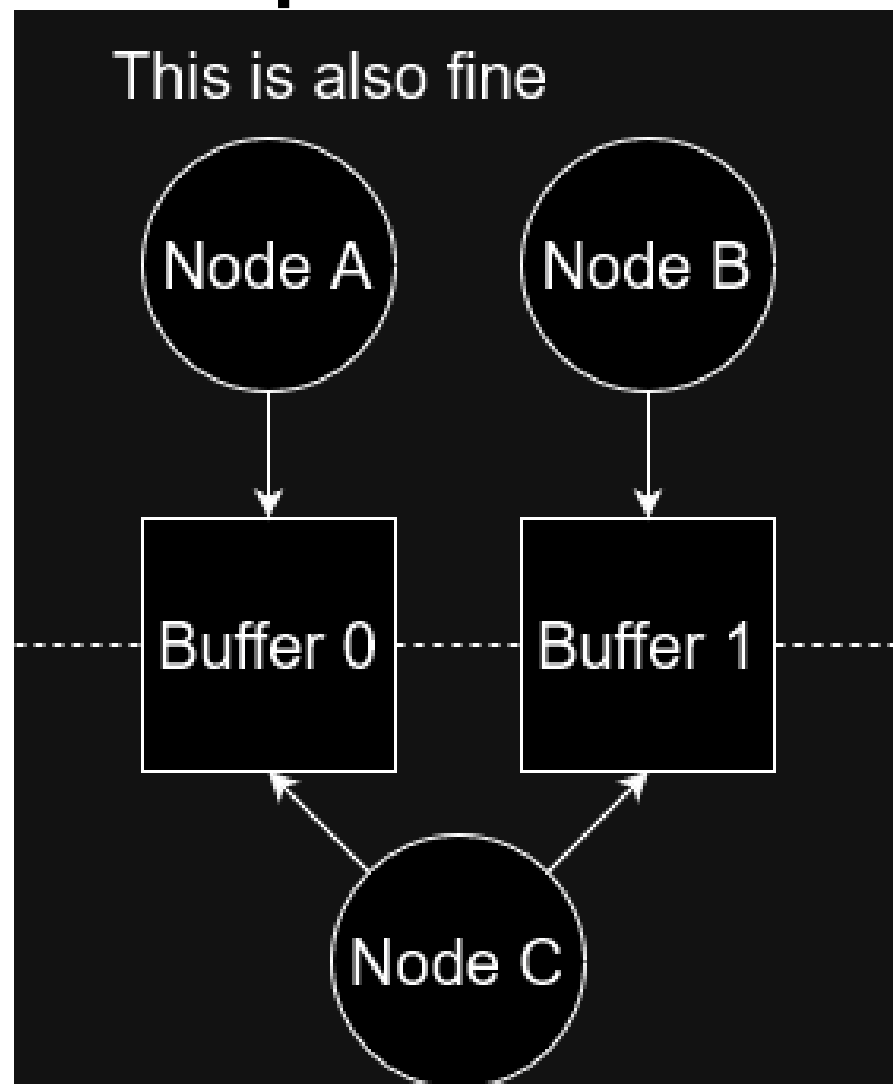
Computational Graph




Computational Graphs



Computational Graphs



Computational Graphs & The Borrow Checker



Computational Graphs

These are the same restrictions enforced by Rust's borrow checker

If we build our own computational graph system, we can enforce our own restrictions and checks

This can happen either at compile time, just-in-time or at runtime

Garbage Collectors

Reference counting in languages like Python, C# and Java

Everything is hidden behind a shared pointer!

Add a garbage collector!

Exercises

[Draw, write and discuss \(the top exercises\)](#)

Once you are done

Do more [Advent of Code](#) or other Rust tutorials if you feel like you need to do more.

or

Do [Ray Tracing In One Weekend](#) – it will confront you with the borrow checker, smart pointers, traits and dyn. I have a code snippet for showing the ray traced image on your screen.

If you complete your 1-to-1 Rust implementation of RTIOW, try to do the following:

- Remove the use of dynamic dispatching by not using dyn anywhere. (Hint: use enums)
- Remove the use of smart pointers. (Hint: my implementation was based on indices and dependency injection of a geometry service)
- Parallelize the application by using the rayon library